

Software Development on Linux Systems

4002-XXX-XX

By

Cody Van De Mark

This work is licensed under the Creative Commons Attribution-ShareAlike 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/3.0/> or send a letter to Creative Commons, 444 Castro Street, Suite 900, Mountain View, California, 94041, USA.

Today

- Integrating with Web Services
- Apache & nginx
- Django

Integrating with Web Services

- GNU/Linux is the most popular platform for web services
- This is because web service support and integration is prevalent on Linux
- Similarly to the application services in the last lecture, integrating with web services on Linux is well supported
- Web service integration on Linux leverages many of same tools for application services, as well as many new ones

Integrating with Web Services

- Web services are supported in a plethora of languages on Linux
- As web services and applications use standard languages (PHP, Python, Java, etc), all integration tools/libraries that exist for those languages can be leveraged even if they are not specific to the web
- There are many extremely well supported frameworks and server engines on Linux;

Integrating with Web Services

- Highly Popular Frameworks for Linux:

Java

Struts

JSF

Spring

Google Web Toolkit

Ruby

Ruby on Rails

Flash

Adobe Flex

Python

Django

Turbogears

CherryPy

Pylons

PHP

CakePHP

CodeIgniter

Symfony

Zephyr

Zend

Drupal (really a CMS)

Integrating with Web Services

- Due to all of the open source languages and integration tools for Linux, it has become a very popular platform
- There are large number of highly supported open source web servers available for Linux:
 - **Apache**
 - **nginx**
 - **Lighttpd**
 - **Hiawatha**
 - **Cherokee**

Integrating with Web Services

- You very often will hear about LAMP:

LAMP - Linux, Apache, MySQL and PHP/Python/Perl

- There are many similar setups available and commonly used
Replacements:

Apache - nginx, Lighttpd, Hiawatha, Cherokee

MySQL - Postgresql, SQLite, CouchDB

PHP - Python, Perl, Ruby, Java

Apache & nginx

- Of the many web servers available for Linux, Apache and nginx have the greatest momentum
- Apache and nginx are currently the two most popular web servers
- Apache itself is the most popularly used and mentioned web server
- nginx is quickly rising as another open source web server, but it is still very new (2004) compared to Apache (1995)

Apache & nginx

- Both Apache and nginx are strongly supported on Linux, with most packages available for both server
- Apache likely has more support overall as it has been around longer, but nginx is mature and mostly has equivalent support
- nginx is an event driven server, while Apache is thread/process driven:

Thread/Process Driven	-	Individual threads for each connection
Event Driven	-	Event fires an action instead of continually checking to see if action occurred

*Note: Events and Threads are not mutually exclusive

Apache

- In lab, you will be working with Apache
- Apache is configured using several configuration files
- Here are some of the main configuration files in Apache
 - apache2.conf** - Top level Apache file, contains global settings
 - httpd.conf** - Main Apache file, contains user specific configuration options
 - ports.conf** - Configure which ports Apache is listening on

Apache

- Apache stores website configurations in a folder called sites-available:

`/etc/apache2/sites-available`

- This folder houses configuration files for each website and their settings
- The configuration files point to a folder where each website resides
- These folders usually contain an index.html file and all of the site pages, css, javascript, images, etc

Apache

- By default, the sites-available folder has a file called default, which is a page you can use to test Apache
- This gives you a starting point for making your own configuration file as it can be referenced for commands and syntax
- The default file points to a pre-made Apache directory that contains the testing pages

Apache

- Once you have a site configured, you can enable it with the command **a2ensite**
- This command stands for “Apache2 Enable Site”
- Once a site is enable, a duplicate of the configuration file from the sites-available folder is made in the folder **sites-enabled**

`/etc/apache2/sites-enabled`

Apache

- Sites are disabled with the command **a2dissite** which stands for “Apache2 Disable Site”
- When this command is run, the configuration file is removed from the sites-enabled folder, and will only reside in sites-available
- If a site is enabled, you may connect to it with the address given in the configuration file
- When connecting from the local machine, the address is given as `http://localhost`

Apache

- At this point, Apache can run, but it will not allow HTTPS connections yet
- To enable a secure connection, you first need a security certificate
- Generating a certificate is done in three steps:
 - 1) Generate RSA Key
 - 2) Generate Certificate Signing Request (CSR)
 - 3) Generate Privacy Enhanced Mail Certificate (PEM)

Apache

- Generating an RSA key is done with openssl using the genrsa flag

```
openssl genrsa -des3 -out mySite.key 2048
```

- Once you have the RSA key, you can generate the certificate signing request using openssl and the req flag

```
openssl req -new -key mySite.key -out mySite.csr
```

- Finally, you can make the certificate with openssl and the x509 flag

```
openssl x509 -req -days 365 -in mySite.csr -signkey mySite.key -out mySite.pem
```


Apache

- After you have generated the certificate, it can be moved into
`/etc/ssl/certs/`
- This is a folder designed to hold all of your certificates
- You will also need to move the RSA key to
`/etc/ssl/private`
- This folder is intended to hold your private keys

Apache

- Apache will automatically reference certificates and keys in `/etc/ssl/certs` and `/etc/ssl/private`
- To enable Apache to do this, you need to use the command **`a2enmod ssl`**
- A2enmod stands for “Apache2 Enable Module”
- This will enable Apache to use SSL certificates and keys

Apache

- Once the SSL module has been enabled, you will find a default HTTPS configuration file in sites-available called **default-ssl**
- If you open the file, you will find that it references Apache's own test key and certificate called:

ssl-cert-snakeoil.pem
ssl-cert-snakeoil.key

- This serves as a testing and reference configuration file for SSL enabled sites

Apache

- When you make your own SSL site configuration files, you need to use your own key and certificate in configuration file
- If you were working with a certificate authority to make the SSL site a legitimate site, you would give your certificate them to validate
- After you finish your SSL site configuration file, you can enable the site as normal using **a2ensite**
- This time, however, you need to restart Apache to load the SSL mod

Apache

- When Apache restarts, you will be able to connect to your site page using an HTTPS connection
- Your browser will give you an untrusted connection error if you have not given the certificate to a valid certificate authority
- The browser will not have any way of validating the certificate unless it can connect to a certificate authority
- You may add the certificate to the browser to see your page

Apache

- You will occasionally see sites that do not use a valid CA and thus their certificate gives an untrusted connection error
- Many times companies do this on their internal network, where the site is not outside facing
- In these instances, they are acting as their own certificate authority and validating the certificate on their own network
- The browser will still not show a legitimate connection as there are no certificate authorities validating the certificate

Django

- Apache is a web server and it controls web connections
- Apache will use the configuration files to use specific folders containing web site files, but it is only hosting static pages
- Apache itself does not generate dynamic pages/content that many sites use
- Instead, this requires software able to generate dynamic content

Django

- Dynamic content is generated one of two ways:
 - Using custom code/engine
 - Using a framework
- Using a framework has more overhead than custom code, but it is much easier and tends to have better security measures in place
- Many of the more popular frameworks are scalable and have existing code to integrate with many different open source technologies

Django

- In Lab, you will be looking at Django, a popular Python framework
- Django is a model-view-controller (MVC) framework that is designed for ease of use, security and scalability
- Django includes many development tools, such as :
 - validation system
 - serialization
 - internationalization
 - unit-testing
 - caching
 - middleware support
 - templates

Django

- Django can work with many open source technologies
- Django has support for various web servers, common gateway interfaces, databases and modules
- One of the nice features Django has, is that development can be done on its own and tested against Django's internal development-only server

Django

- Django's internal development-only server can not be used for hosting live pages, as it is not secure and not designed for performance
- The internal development server allows you to thoroughly test pages as if they were live before deploying to a web server
- This allows you to do web development and testing before building a server environment or deploying to an existing environment

Django

- We are using Django in lab because:
 - It is one of the easiest frameworks to use,
 - It is very secure
 - It is one of the fastest
 - It illustrates how well open source technologies integrate together
- Django has tremendous integration support for many technologies out of the box
- In lab, you will look at integrating Django with MySQL

Django

- The first step in setting up Django, is to choose your database
- In the database, you need to create a user for Django to use and create a database named “Django”

*Note: your database user must have access to the Django table

- Django will use this database to store all of its data and user data

*Note: The database does not necessarily need to be a relational database; Many NoSQL databases are supported as well

Django

- Once the database has been configured, you need to create a folder to house your Django project
- Inside of that folder, you can initiate Django with
`django-admin startproject myProjectName`
- This will create a directory called myProjectName filled with Django configuration files

Django

- In the new folder, there will be a file called settings.py
- This is a global setting configuration file for your Django project
- In the file there is an 'ENGINE' variable set to `django.db.backends`.
- You can interact with many databases by tacking on the name of the database to the end of the variable

Example: `django.db.backends.mysql`

Django

- Below the engine line are variables for the database name, database username and database user password
- Once these have been filled in and the file has been saved, Django can setup the database using the manage script in the same folder

`./manage.py syncdb`

- Once the database is setup, you can test Django using Django's internal development server

`./manage.py runserver`

Django

- Once the development server is running, you should be able to connect to it on your local machine on port 8000
- If you go to <https://127.0.0.1:8000> in a web browser, a web page should load saying “It Worked”
- After Django is configured, you can start creating apps with
`./manage.py startapp myAppName`

Django

- After startapp has been run, there will be a folder called myAppName that is full of application configuration files
- After the app has been made, you can open the settings.py file again and find the INSTALLED_APPS variable;

At the end of the variable you can add your new application name
'myAppName'

- This allows Django to know which apps you want to be running

Django

- In Django, content is generated dynamically
- When a request is made, Django calls a function to generate the page you want them to have based on their specific data
- You will notice in the main Django folder, there is a file called “urls.py”
- This file directs all urls to the functions you want them to call based on regular expressions (regex)

Django

- You will also find that there is a `urls.py` file inside of the app folder
- This is because the first `urls.py` is Django's global level url management;

This file determines which app to call based on the url

- After the app is called, the app takes the url into its own `urls.py` file to determine which page to generate

Django

- An example of the global level urls regex

```
url(r'^social/', include('socialNetwork.urls'))
```

This line would take any urls that begin with “social” after the address and call the application called “socialNetwork”

- An example of the application level urls regex

```
(r'^login/$', 'createLoginPage')
```

This line would take any urls that end in “login/” and call the function createLoginPage to generate the page

Django

- These functions that generate pages are in the file **views.py**
- The `views.py` file contains functions for each url call in order to generate any data necessary
- Most of the time pages are generated using html pages that have been created at templates
- The views function can call these templates and fill in any customized data the page needs

Django

- Template html pages can be written like normal html pages, but can contain variables and evaluation statements that Django can alter
- Variables in Django html pages are made with

`{{ variableName }}`

- Django can then grab that variable name in a function and alter it to the value it desires in proper html format

Django

- Similarly, Django can use evaluation statements in html pages
- Evaluation statements are done with

```
{% for v in variableName %}
```

```
{% endfor %}
```

- This allows html and variables in between the statements to be generated as needed by the loop;

Django allows many different types of evaluation to be done

Django

- With the powerful features Django offers and its strong integration with open source, dynamic web pages can be made quickly
- Django easily integrates with many open source technologies with little configuration or extras to install
- As Django itself is open source, new developments and integration libraries are being made often